


Lecture 13

RNN






Sequence data


- We don't understand one word only
- We understand based on the previous words + this word. (time series)
- NN/CNN cannot do this



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Example

Speech recognition		→	"The quick brown fox jumped over the lazy dog."
Music generation		→	
Sentiment classification	"There is nothing to like in this movie."	→	
DNA sequence analysis	AGCCCCTGTGAGGAACTAG	→	AGCCCCTGTGAGGAACTAG
Machine translation	Voulez-vous chanter avec moi?	→	Do you want to sing with me?
Video activity recognition		→	Running



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Recurrent Neural Network

We can process a sequence of vectors \mathbf{x} by applying a recurrence formula at every time step:


$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state

some function with parameters W

old state

input vector at some time step



Recurrent Neural Network


We can process a sequence of vectors \mathbf{x} by applying a recurrence formula at every time step:

$$h_t = f_W(h_{t-1}, x_t)$$

Notice: the same function and the same set of parameters are used at every time step.

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$


$$y_t = W_{hy}h_t$$



Character-level language model example

Vocabulary:
[h,e,l,o]


Example training
sequence:
“hello”



Character-level language model example

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”

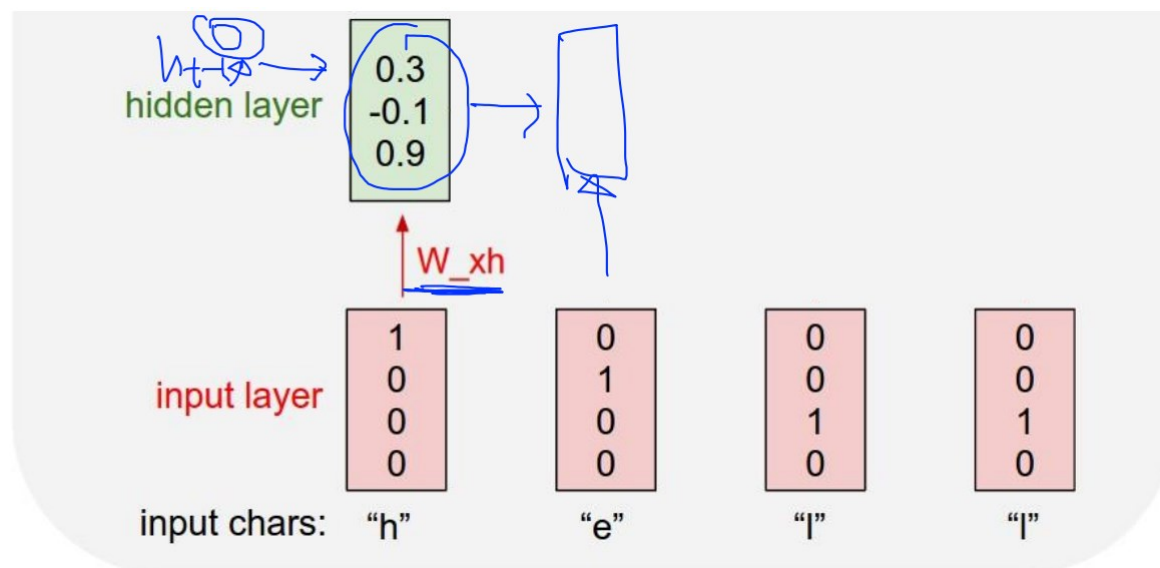


Character-level language model example

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

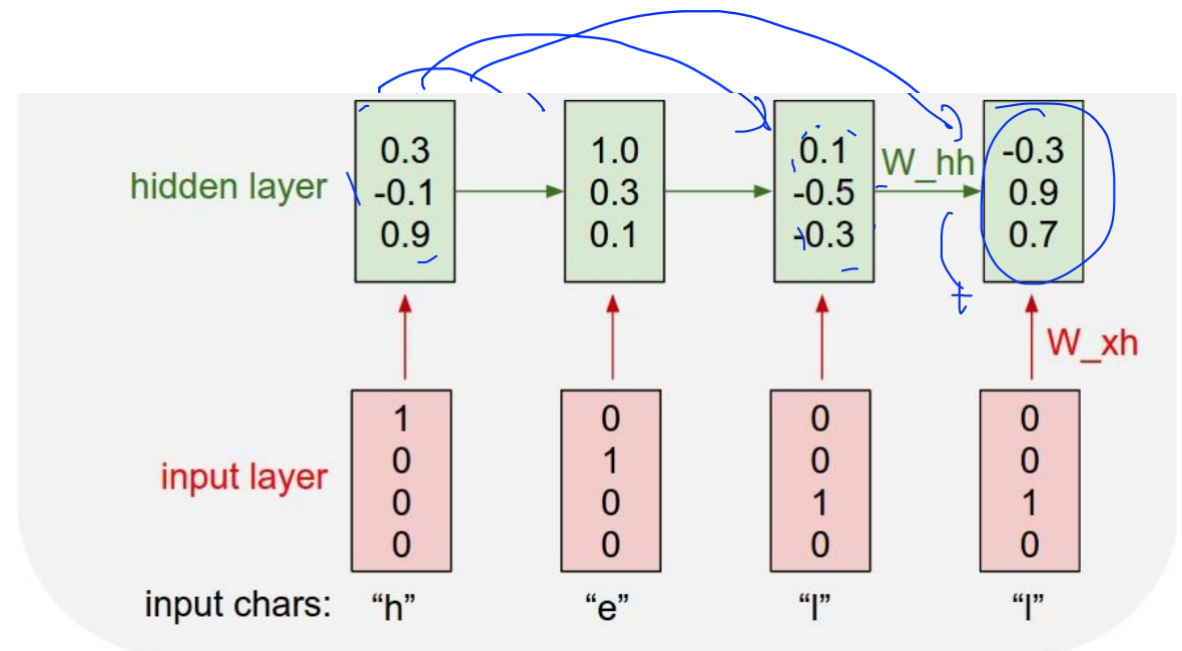


Character-level language model example

Vocabulary:
[h,e,l,o]

Example training sequence:
“hello”

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

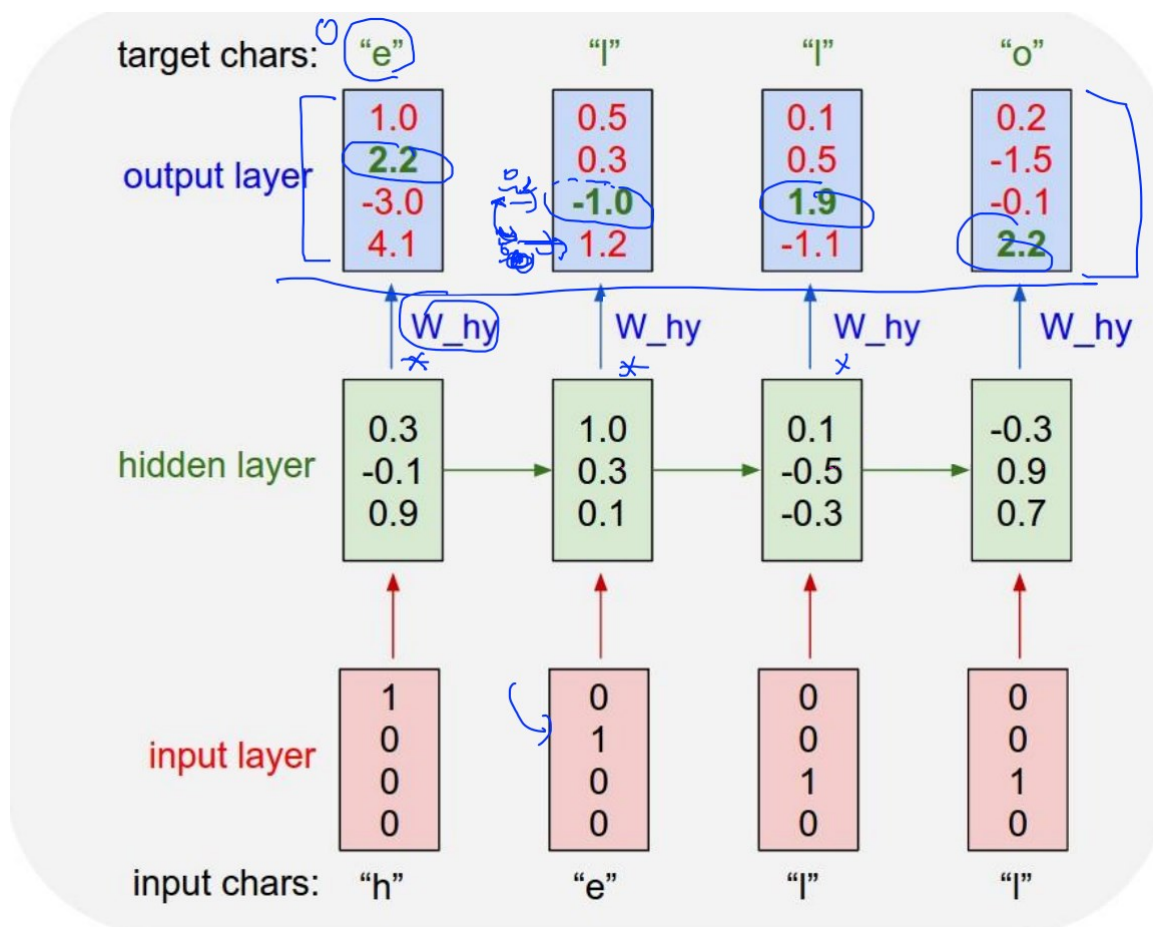


Character-level language model example

$$y_t = W_{hy}h_t$$

Vocabulary:
[h,e,l,o]

Example training
sequence:
“hello”




RNN applications


- https://github.com/TensorFlowKR/awesome_tensorflow_implementations
 - Language Modeling
 - Speech Recognition
 - Machine Translation
 - Conversation Modeling/Question Answering
 - Image/Video Captioning
 - Image/Music/Dance Generation

Recurrent Networks offer a lot of flexibility:


one to one




one to many




many to one



many to many



many to many



**Vanilla
Neural
Networks**

e.g. **Image
Captioning**
image ->
sequence of
words


e.g. **Sentiment
Classification**
sequence of
words ->
sentiment

e.g. **Machine
Translation**
seq of words ->
seq of words


e.g. **Video
classification
on frame level**

Understanding LSTM Networks

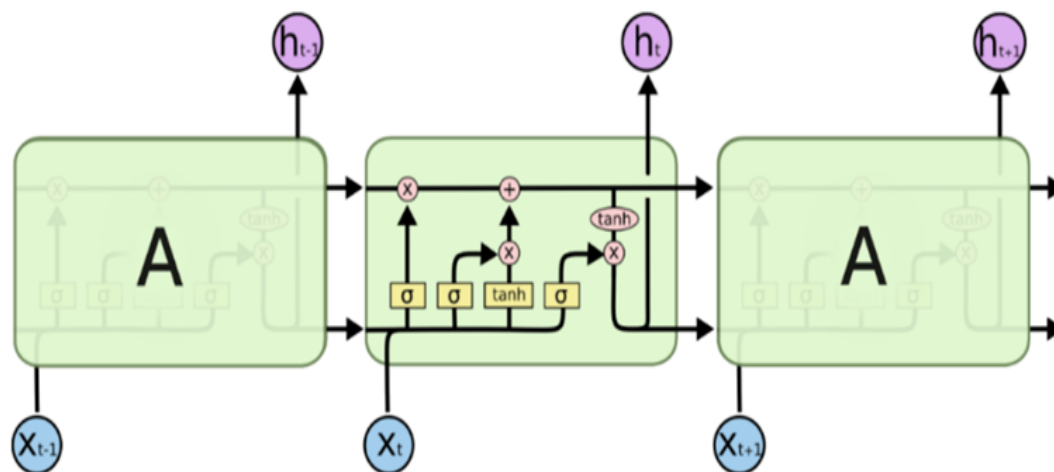
- <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>



RNN vs LSTM




The repeating module in a standard RNN contains a single layer.



The repeating module in an LSTM contains four interacting layers.


Core idea behind LSTM

- Key to LSTMs is the **cell state**
 - The horizontal line running through the top of the diagram
- LSTM can add or remove information to the cell state
- How? Through regulated structures called **gates**.
- LSTM has **three gates** to **protect** and **control** cell state



Steps


- **Forget gate layer** decides what information will be thrown away
- Looks at h_{t-1} and x_t and outputs a number between 0 and 1
- 1 represents **completely keep this**, 0 represents **completely get rid of this**
- **Example**: forget the gender of the old subject, when we see a new subject



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$


Steps

- **Next step:** decides what new information will be stored in the cell state
- Two parts –
 - A **sigmoid** layer (**input gate layer**): decides what values we'll update
 - A **tanh** layer: creates a vector of new candidate values, \tilde{C}_t
- **Example:** add the gender of the new subject to the cell state
 - Replace the old one we're forgetting




Steps

- **Next step:** update old state by C_{t-1} into the new cell state C_t
- Multiply old state by f_t
 - Forgetting the things we decided to forget earlier
- Then we add $i_t * \tilde{C}_t$



Steps

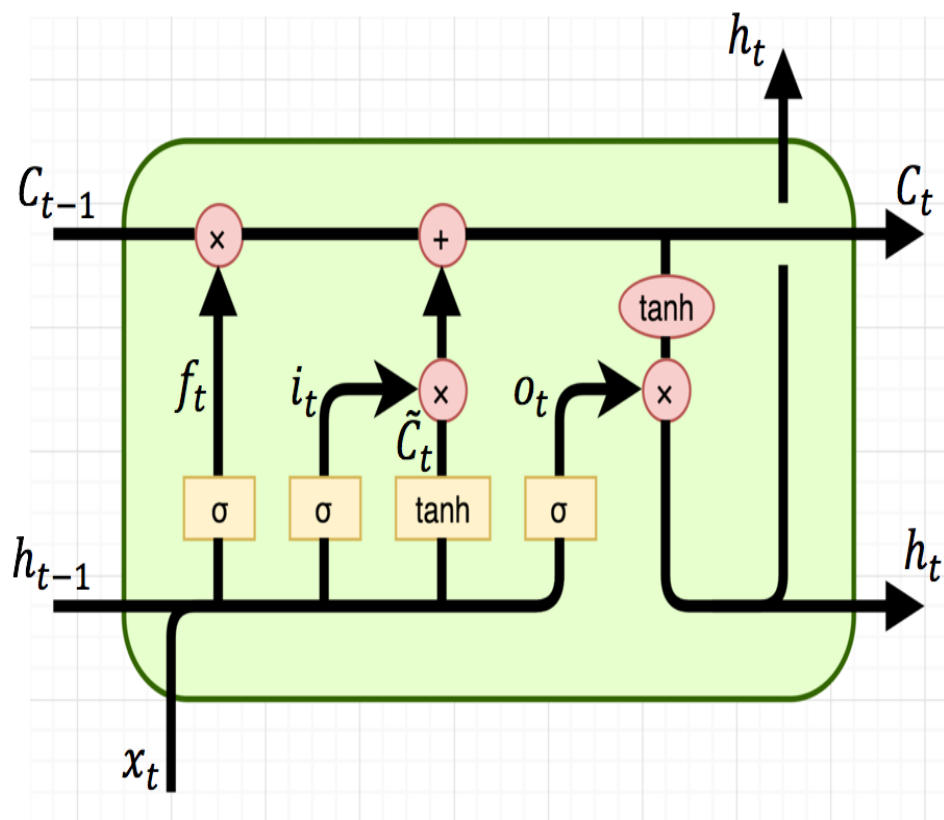
- **Final step:** decide what we're going to output
- First, we run a **sigmoid layer**
 - Which decides what parts of the cell state we're going to output
- Then, we put the cell state through **tanh** and multiply it by the output of the sigmoid gate




$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$

LSTM vs GRU



(a) Long Short-Term Memory



(b) Gated Recurrent Unit

Conclusion



- Vanilla RNNs are simple but don't work very well
- LSTM is a big step in what we can accomplish with RNNs
 - Backward flow of gradients in RNN can explode or vanish. Exploding is controlled by gradient clipping and vanishing is controlled with additive interactions (LSTM)
- GRU's are faster to train and need fewer data to generalize.
- When there is enough data, an LSTM's greater expressive power may lead to better results
- Better understanding (both theoretical and empirical) is needed